

SELECT ROUND-UP

Table of Contents

- [README](#)
- [Download](#)
- [Preliminaries](#)
- [Create some tables](#)
- [Simple JOINS](#)
- [JOIN ... ON](#)
- [WHERE examples](#)
- [GROUP BY Examples](#)
- [ORDER BY examples](#)
- [LIMIT and OFFSET examples](#)

README

This notebook contains a round-up of the entire SELECT command pipeline in SQL. The code blocks run with SQLite, and we use a simple database of three different tables to explore SELECT.

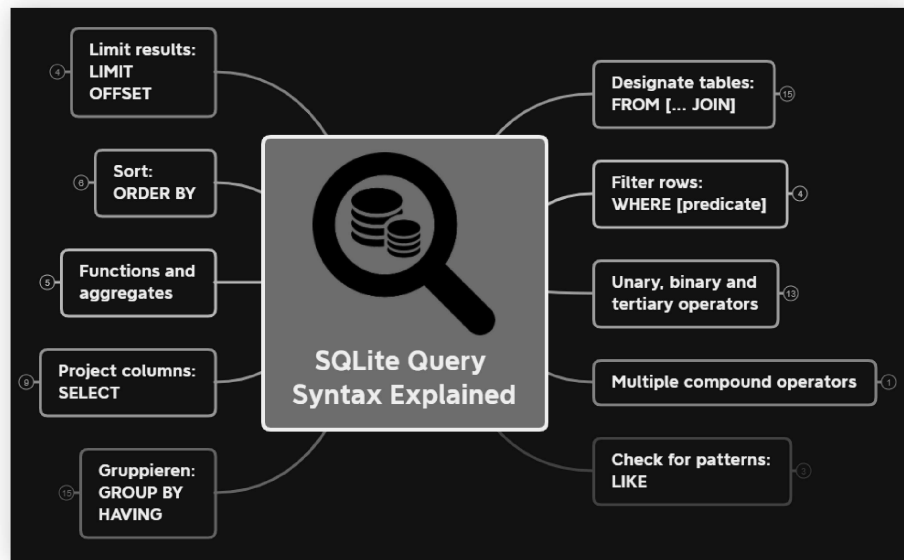
To run this notebook, which resides in GitHub and GDrive (in the *practice directory*), you need to have *SQLite* installed and in the *PATH* so that Emacs can find it, and you need to have the `~.emacs~` file also available from GDrive, installed in Emacs' HOME directory.

Download

- Download `xyz.zip` [from GDrive](#).
- This includes the SQLite file `xyz.sql` and a directory with image files.

Preliminaries

- For the presentation, I'm going to use the XMind-map from GitHub.



Presented with XMind

Figure 1: SELECT round-up

- [X]

How are the commands in the mindmap ordered (from right to left)?

Answer: this is the pipeline order for SELECT

Create some tables

- [X]

Let's load some tables to play with. The file `xyz.sql` is available from [GDrive's notebook directory](#). It contains a complete database dump. We load it with the SQLite command `.read`, then check database and tables. There should be three, named `x`, `y`, and `z`.

```
.read xyz.sql
.database
.tables
```

```
main: c:\Users\birkenkrahe\Documents\GitHub\db330\practice\xyz.db r/w
x y z
```

- []

Challenge: if the tables already exist in `xyz.db` (and you don't have the `IF NOT EXISTS` clause in the `CREATE TABLE` statement), you need to delete the database before loading the tables. But even if you do that, the `INSERT` commands will add to the existing tables. So how can you limit the import or delete duplicate entries?

ANSWER: You have to select the unique entries with `DISTINCT`.

- We use `SELECT` with the wildcard `*` for the column selection, and without conditions (i.e. conditions on row selections) to look at each table.
- []

Print table `x` - one integer column and one column of strings.

```
SELECT * FROM x;
```

	a	b
	-	-----
1		Alice
2		Bob
3		Charlie

- []

Print table `y` - one integer column and one column of floats.

```
SELECT * FROM y;
```

	c	d
	-	-----
1		3.14159
1		2.71828
2		1.61803

- []

Print table `y` - two integer columns

```
SELECT * FROM z;
```

	a	e
	-	---
1		100
1		150
3		300
9		900

<pre>sqlite> SELECT * FROM x; a b --- --- 1 Alice 2 Bob 3 Charlie</pre>	<pre>sqlite> SELECT * FROM y; c d --- --- 1 3.14159 1 2.71828 2 1.61803</pre>	<pre>sqlite> SELECT * FROM z; a e --- --- 1 100 1 150 3 300 9 900_</pre>
---	---	---

Figure 2: Sample tables x,y,z

Simple JOINS

- []

CROSS JOIN: Since both tables had 3 rows and 2 columns, the result set has $9=3*3$ rows and $4=2*2$ columns.

```
SELECT * FROM x JOIN y;
```

a	b	c	d
-	-----	-	-----
1	Alice	1	3.14159
1	Alice	1	2.71828
1	Alice	2	1.61803
2	Bob	1	3.14159
2	Bob	1	2.71828
2	Bob	2	1.61803
3	Charlie	1	3.14159
3	Charlie	1	2.71828
3	Charlie	2	1.61803

```
SELECT * FROM x CROSS JOIN y;
```

a	b	c	d
-	-----	-	-----
1	Alice	1	3.14159
1	Alice	1	2.71828
1	Alice	2	1.61803
2	Bob	1	3.14159
2	Bob	1	2.71828

2	Bob	2	1.61803
3	Charlie	1	3.14159
3	Charlie	1	2.71828
3	Charlie	2	1.61803

```
SELECT * FROM x,y;
```

a	b	c	d
-	-----	-	-----
1	Alice	1	3.14159
1	Alice	1	2.71828
1	Alice	2	1.61803
2	Bob	1	3.14159
2	Bob	1	2.71828
2	Bob	2	1.61803
3	Charlie	1	3.14159
3	Charlie	1	2.71828
3	Charlie	2	1.61803

JOIN ... ON

- []

INNER JOIN: Remember the zipper principle - identify 2 columns to zip xtogether. This identification follows after the ON keyword.

```
SELECT * FROM x JOIN y ON a = c;
```

a	b	c	d
-	-----	-	-----
1	Alice	1	2.71828
1	Alice	1	3.14159
2	Bob	2	1.61803

- o Compare with the CROSS JOIN before. This time, Only those columns that satisfy the condition a=c are included in 1

- []

What if we want to JOIN tables x and z? They both have a column named a. We now need to qualify the selection with .

```
SELECT * FROM x JOIN z ON x.a = z.a;
```

a	b	a	e
---	---	---	---

```

-  - - - - -  -  - - -
1  Alice    1  100
1  Alice    1  150
3  Charlie  3  300

```

- []

You can resolve column name confusions (same column name in different tables) by using aliases. Run the previous command `1` again, but alias `x` as `t_x` and `z` as `t_z`.

```
SELECT * FROM x AS t_x JOIN z AS t_z ON t_x.a = t_z.a;
```

```

a  b      a      e
-  - - - - -  -  - - -
1  Alice    1  100
1  Alice    1  150
3  Charlie  3  300

```

- []

There are five (!) more JOIN commands:

- LEFT OUTER JOIN will also include not matched items
- COMPOUND JOIN joins multiple tables

```
SELECT * FROM x JOIN y ON x.a=y.c LEFT OUTER JOIN z ON y.c=z.a;
```

```

a  b      c      d      a      e
-  - - - - -  -  - - - - -  - - - - -  - - - - -
1  Alice  1  2.71828      1      100
1  Alice  1  2.71828      1      150
1  Alice  1  3.14159      1      100
1  Alice  1  3.14159      1      150
2  Bob    2  1.61803  [NULL]  [NULL]

```

- Work through this example until you really understand what this multiple join = INNER JOIN + LEFT JOIN does!

WHERE examples

- []

Print a row: the value Alice for the attribute `x.b`.

```
SELECT * FROM x WHERE x.b = 'Alice';
```

```

a  b
-  - - - - -

```

1 Alice

- []

Print a range of values, for $1.0 < d < 3.0$.

```
SELECT * FROM y WHERE y.d BETWEEN 1.0 AND 3.0;
```

c d

- -----

1 2.71828

2 1.61803

- []

Print columns c, d and a column for the sum of c+d with the condition that the sum is smaller than 4.

```
SELECT c, d, c+d AS sum FROM y WHERE sum < 4.0
```

c d sum

- -----

1 2.71828 3.71828

2 1.61803 3.61803

- []

The next block 1 uses foods.db to select a range of values with wildcards. Here, * instead of % would also work (try it).

```
Select name from foods where name between 'Ta%' AND 'Ti%';
```

Tarragon

Tea

Three Musketeers

Tamale

Tamales

GROUP BY Examples

- []

Group table z by the column z.a. Can you guess how many rows are going to be printed?

```
SELECT a FROM z GROUP BY z.a;
```

a

-

1

3

9

- []

Print the number of rows next to every value of z.a. Call this new column 'count'n.

```
SELECT a, COUNT(a) AS count FROM z GROUP BY z.a;
```

a	count
1	2
3	1
9	1

-

1 2

3 1

9 1

- []

Run 1 again (group by z.a) but now also print out the sum of all the z.e values in each group. Call the new column 'TOTAL'.

```
SELECT a, sum(e) AS TOTAL FROM z GROUP BY z.a;
```

a	TOTAL
1	250
3	300
9	900

-

1 250

3 300

9 900

- []

Run 1 again (group by z.a) but now also compute

- the sum(e) as SUM
- the count(e) as TOTAL
- the average as AGG computed with sum and count
- the average as AVG computed with the aggregate function

```
SELECT a, sum(e) as SUM, count(e) as TOTAL, sum(e)/count(e) AS AGG, avg(e) AS AVG FROM z GROUP BY z.a;
```

a	SUM	TOTAL	AGG	AVG
1	250	2	125	125.0
3	300	1	300	300.0
9	900	1	900	900.0

-

1 250 2 125 125.0

3 300 1 300 300.0

9 900 1 900 900.0

- []

A HAVING clause can be used to filter rows based off the results of the sum() aggregation. Run the block 1.

```
SELECT a, sum(e) AS TOTAL FROM z GROUP BY z.a HAVING total > 500;
```

a	TOTAL
---	-------

-	-----
---	-------

9	900
---	-----

- []

An example with the foods database, and the table foods; print the food type ID and the total number of food types per food type group, and print those IDs whose group has less than 20 foods in it.

```
SELECT type_id, COUNT(*)
FROM foods
GROUP BY type_id
HAVING COUNT(*) < 20;
```

type_id	COUNT(*)
---------	----------

-----	-----
-------	-------

2	15
---	----

5	17
---	----

6	4
---	---

11	16
----	----

13	14
----	----

14	19
----	----

type_id	COUNT(*)
---------	----------

-----	-----
-------	-------

2	15
---	----

5	17
---	----

6	4
---	---

11	16
----	----

13	14
----	----

14	19
----	----

ORDER BY examples

- []

Order table y by the numbers in y.d, and print all columns.

```
SELECT * FROM y ORDER BY d;
```

c	d
---	---

```
- -----
2 1.61803
1 2.71828
1 3.14159
-----
c      d
```

```
- -----
2 1.61803
1 2.71828
1 3.14159
-----
```

- []

An example from the foods database. Display all columns from the table `foods`, and filter those rows whose name begins with a B. Print only 10 lines.

```
SELECT * FROM foods WHERE name LIKE 'B%' LIMIT 10;
```

```
id  type_id  name
--  -
1   1        Bagels
2   1        Bagels, raisin
3   1        Bavarian Cream Pie
4   1        Bear Claws
5   1        Black and White cookies
6   1        Bread (with nuts)
7   1        Butterfingers
48  2        Bran
63  3        Broiled Chicken
87  4        Barbeque Sauce
```

- []

Now take the command from 1 and order by food type ID in descending order. To do this, put `DESC` after the column name.

```
SELECT * FROM foods WHERE name LIKE 'B%'
ORDER BY type_id DESC
LIMIT 10;
```

```
id  type_id  name
---  -
382  15       Baked Beans
383  15       Baked Potato w/Sour Cream
384  15       Big Salad
385  15       Broccoli
362  14       Bouillabaisse
326  12       Bologna
327  12       Bacon Club (no turkey)
328  12       BLT
329  12       Brisket Sandwich
274  10       Bacon
```

- []

You can order by any number of columns. Modify the command in 1: order in descending order on type_id, and then order the result in ascending order alphabetically on name.

```
SELECT * FROM foods WHERE name LIKE 'B%'
ORDER BY type_id DESC, name ASC
LIMIT 10;
```

id	type_id	name
---	-----	-----
382	15	Baked Beans
383	15	Baked Potato w/Sour Cream
384	15	Big Salad
385	15	Broccoli
362	14	Bouillabaisse
328	12	BLT
327	12	Bacon Club (no turkey)
326	12	Bologna
329	12	Brisket Sandwich
274	10	Bacon

LIMIT and OFFSET examples

- []

From foods, print the first 5 lines of all columns.

```
SELECT * FROM foods LIMIT 5;
```

id	type_id	name
--	-----	-----
1	1	Bagels
2	1	Bagels, raisin
3	1	Bavarian Cream Pie
4	1	Bear Claws
5	1	Black and White cookies

- []

Change the code in 1 only to skip the first 2 lines.

```
SELECT * FROM foods LIMIT 5 OFFSET 2;
```

id	type_id	name
--	-----	-----
3	1	Bavarian Cream Pie
4	1	Bear Claws
5	1	Black and White cookies
6	1	Bread (with nuts)
7	1	Butterfingers

- []

Skip the first 2 lines without using the keyword OFFSET.

```
SELECT * FROM foods LIMIT 2, 5;
```

id	type_id	name
--	-----	-----
3	1	Bavarian Cream Pie
4	1	Bear Claws
5	1	Black and White cookies
6	1	Bread (with nuts)
7	1	Butterfingers

Author: Marcus Birkenkrahe

Created: 2022-03-11 Fri 22:32

[Validate](#)